

	Number of Concurrent Requests made ->					
Metrics V			1	10	50	100
		Average TTFT	25 ms	30 ms	52 ms	65 ms
	Latency	Average ITL	8.4 ms	9 ms	13 ms	18 ms
	Throughput	Generation Token Throughput Peak Value	120 tokens/s	1100 tokens/s	3300 tokens/s	4000 tokens/s
	Memory	Peak Cache Utilization	0.26 %	0.85 %	3.25 %	5.9 %
	Number of Requests made = 3000					
	Request Rate = 20 r/s					
	GPU : A10					
	PCIe					