

	Category	Metric					
	Model	Model : Qwen2.5-1.5B-Instruct					
		Number of Prompts : 3000					
	vLLM server	Request Rate : 20 requests/s					
		GPU : 1 x A10					
		Interconnect : PCIe					
	System	Parallelism: None					
	Metrics vs Number of Concurrent Requests						
	Number of Concurrent Requests ->						
Metrics V			1	10	50	100	500
		Mean TTFT	27 ms	30 ms	52 ms	65 ms	60 ms
	Latency	Mean ITL	8.4 ms	9 ms	13 ms	18 ms	18 ms
	Throughput	Generation Token Throughput Peak Value	117 tokens/s	1100 tokens/s	3300 tokens/s	4000 tokens/s	4000 tokens/s
	Memory	Peak Cache Utilization	0.26 %	0.85 %	3.25 %	5.9 %	5.1 %
	Observations Table						
	Observation ->						
Metrics V		1 to 100 Concurrent Requests	At 500 Concurrent Requests (comparison with 100 Concurrent Requests)				
	Mean TTFT	Steady Climb	5 ms drop (minor)				
	Mean ITL	Steady Climb	Remains approximately the same				
	Generation Token Throughput Peak Value	Steady Climb	Remains approximately the same				
	Peak Cache Utilization	Steady Climb	8 % drop (minor)				