

	Category	Metric					
	Model	Qwen2.5-1.5B-Instruct					
	vLLM server	Number of Prompts : 3000					
		GPU : 1 x A10					
	System	Parallelism: None					
	Max Concurrent Requests ->						
Metrics V			1	10	50	100	500
		Mean TTFT	21.82 ms	36.15 ms	59.71 ms	106.10 ms	9835.48 ms
	Latency	Mean ITL	8.54 ms	9.28 ms	14.40 ms	23.04 ms	56.04 ms
	Throughput	Generation Token Throughput Peak Value	121 tokens/s	1150 tokens/s	4080 tokens/s	5760 tokens/s	8767 tokens/s
	Memory	Peak Cache Utilization	0.28%	0.86%	3.44%	6.21%	15.10%
	Observation ->						
Metrics V		1 to 100 Concurrent Requests	At 500 Concurrent Requests				
	Mean TTFT	Nearly Doubles each step	Large increase by ~100x				
	Mean ITL	Increase	Increase				
	Generation Token Throughput Peak Value	Increase	Increase				
	Peak Cache Utilization	Increase	Increase				