

	Category	Metric					
	Model	Model : Qwen2.5-1.5B-Instruct					
	vLLM server	Number of Prompts : 3000					
		GPU : 1 x A10					
	System	Parallelism: None					
	Metrics vs Max Concurrent Requests						
	Max Concurrent Requests ->						
Metrics V			1	10	50	100	500
		Mean TTFT	21.82 ms	36.15 ms	59.71 ms	106.10 ms	9835.48 ms
	Latency	Mean ITL	8.54 ms	9.28 ms	14.40 ms	23.04 ms	56.04 ms
	Throughput	Generation Token Throughput Peak Value	121 tokens/s	1150 tokens/s	4080 tokens/s	5760 tokens/s	8767 tokens/s
	Memory	Peak Cache Utilization	0.28%	0.86%	3.44%	6.21%	15.10%
	Observations Table						
	Observation ->						
Metrics V		1 to 100 Concurrent Requests	At 500 Concurrent Requests				
	Mean TTFT	Nearly Doubles each step	Explodes to ~100x				
	Mean ITL	Climbing	Still Climbing				
	Generation Token Throughput Peak Value	Climbing	Still Climbing				
	Peak Cache Utilization	Climbing	Still Climbing				